

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import os

from sklearn import svm, tree, linear_model, neighbors, naive_bayes, ensemble, discriminant_analysis, g
aussian_process
import tensorflow as tf
from tensorflow.keras

from sklearn import metrics
from sklearn.compose import make_column_transformer

from sklearn.model_selection import cross_val_score
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.model_selection import ShuffleSplit
from sklearn import model_selection

from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import f_classif

In [3]: os.getcwd()

Out[3]: 'C:\Users\robtu\

In [2]: os.chdir(r"C:\Users\robtu\Kaggle Competitions")

In [3]: train_df = pd.read_csv("train.csv")
test_df = pd.read_csv("test.csv")
train_df.head()

Out[3]:
```

	id	cat0	cat1	cat2	cat3	cat4	cat5	cat6	cat7	cat8	...	cont2	cont3	cont4	cont5	cont6	cont7	cont8
0	0	A	I	A	B	B	BI	A	S	Q	...	0.759439	0.795549	0.881917	0.621672	0.592184	0.791921	0.815254
1	1	A	I	A	B	E	BI	K	W	AD	...	0.386385	0.541366	0.386982	0.357778	0.600044	0.408701	0.993553
2	2	A	K	A	A	E	BI	A	E	BM	...	0.343255	0.616352	0.793687	0.552877	0.352113	0.388835	0.412303
3	3	A	A	G	E	B	BI	A	Y	AD	...	0.831147	0.807807	0.800032	0.619147	0.221789	0.897617	0.633669
4	4	A	I	G	B	E	BI	C	G	Q	...	0.338818	0.277308	0.610578	0.128291	0.578764	0.279167	0.351103

5 rows × 32 columns

```
In [4]: df = pd.concat([train_df.drop(columns=['target']), test_df], ignore_index = True)
y_train = train_df['target']
test_ids = test_df['id']
```

## Data Exploration

```
In [5]: train_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300000 entries, 0 to 299999
Data columns (total 32 columns):
#   Column      Non-Null Count  Dtype
---  --
0   id           300000 non-null  int64
1   cat0        300000 non-null  object
2   cat1        300000 non-null  object
3   cat2        300000 non-null  object
4   cat3        300000 non-null  object
5   cat4        300000 non-null  object
6   cat5        300000 non-null  object
7   cat6        300000 non-null  object
8   cat7        300000 non-null  object
9   cat8        300000 non-null  object
10  cat9        300000 non-null  object
11  cat10       300000 non-null  object
12  cat11       300000 non-null  object
13  cat12       300000 non-null  object
14  cat13       300000 non-null  object
15  cat14       300000 non-null  object
16  cat15       300000 non-null  object
17  cat16       300000 non-null  object
18  cat17       300000 non-null  object
19  cat18       300000 non-null  object
20  cont0       300000 non-null  float64
21  cont1       300000 non-null  float64
22  cont2       300000 non-null  float64
23  cont3       300000 non-null  float64
24  cont4       300000 non-null  float64
25  cont5       300000 non-null  float64
26  cont6       300000 non-null  float64
27  cont7       300000 non-null  float64
28  cont8       300000 non-null  float64
29  cont9       300000 non-null  float64
30  cont10      300000 non-null  float64
31  target      300000 non-null  int64
dtypes: float64(11), int64(2), object(19)
memory usage: 73.2+ MB

In [29]: df.shape

Out[29]: (500000, 31)

In [34]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500000 entries, 0 to 499999
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  --
0   id           500000 non-null  int64
1   cat0        500000 non-null  object
2   cat1        500000 non-null  object
3   cat2        500000 non-null  object
4   cat3        500000 non-null  object
5   cat4        500000 non-null  object
6   cat5        500000 non-null  object
7   cat6        500000 non-null  object
8   cat7        500000 non-null  object
9   cat8        500000 non-null  object
10  cat9        500000 non-null  object
11  cat10       500000 non-null  object
12  cat11       500000 non-null  object
13  cat12       500000 non-null  object
14  cat13       500000 non-null  object
15  cat14       500000 non-null  object
16  cat15       500000 non-null  object
17  cat16       500000 non-null  object
18  cat17       500000 non-null  object
19  cat18       500000 non-null  object
20  cont0       500000 non-null  float64
21  cont1       500000 non-null  float64
22  cont2       500000 non-null  float64
23  cont3       500000 non-null  float64
24  cont4       500000 non-null  float64
25  cont5       500000 non-null  float64
26  cont6       500000 non-null  float64
27  cont7       500000 non-null  float64
28  cont8       500000 non-null  float64
29  cont9       500000 non-null  float64
30  cont10      500000 non-null  float64
dtypes: float64(11), int64(1), object(19)
memory usage: 118.3+ MB

In [ ]: y_train.shape

In [6]: train_df['cat0'].value_counts

Out[6]:
```

	count	A
1	A	
2	A	
3	A	
4	A	
...	...	
299995	A	
299996	A	
299997	A	
299998	B	
299999	A	

Name: cat0, Length: 300000, dtype: object>

```
In [7]: test_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  --
0   id           200000 non-null  int64
1   cat0        200000 non-null  object
2   cat1        200000 non-null  object
3   cat2        200000 non-null  object
4   cat3        200000 non-null  object
5   cat4        200000 non-null  object
6   cat5        200000 non-null  object
7   cat6        200000 non-null  object
8   cat7        200000 non-null  object
9   cat8        200000 non-null  object
10  cat9        200000 non-null  object
11  cat10       200000 non-null  object
12  cat11       200000 non-null  object
13  cat12       200000 non-null  object
14  cat13       200000 non-null  object
15  cat14       200000 non-null  object
16  cat15       200000 non-null  object
17  cat16       200000 non-null  object
18  cat17       200000 non-null  object
19  cat18       200000 non-null  object
20  cont0       200000 non-null  float64
21  cont1       200000 non-null  float64
22  cont2       200000 non-null  float64
23  cont3       200000 non-null  float64
24  cont4       200000 non-null  float64
25  cont5       200000 non-null  float64
26  cont6       200000 non-null  float64
27  cont7       200000 non-null  float64
28  cont8       200000 non-null  float64
29  cont9       200000 non-null  float64
30  cont10      200000 non-null  float64
dtypes: float64(11), int64(1), object(19)
memory usage: 47.3+ MB
```

## Feature Engineering

```
In [28]: df_dummies = pd.get_dummies(df).drop(columns='id')
df_dummies.head()

Out[28]:
```

	cont0	cont1	cont2	cont3	cont4	cont5	cont6	cont7	cont8	cont9	...	cat16_C	cat16_D	cat17_A	...
0	0.629858	0.855349	0.759439	0.795549	0.881917	0.621672	0.592184	0.791921	0.815254	0.965006	...	0	1	0	0
1	0.370727	0.328929	0.386385	0.541366	0.388982	0.357778	0.600044	0.408701	0.399353	0.927406	...	0	0	1	0
2	0.502272	0.322749	0.343255	0.616352	0.793687	0.552877	0.352113	0.388835	0.412303	0.292696	...	0	1	0	0
3	0.934242	0.707663	0.831147	0.807807	0.800032	0.619147	0.221789	0.897617	0.633669	0.760318	...	0	1	0	0
4	0.254427	0.274514	0.338818	0.277308	0.610578	0.128291	0.578764	0.279167	0.351103	0.357084	...	0	0	0	0

5 rows × 642 columns

```
In [29]: X_train = df_dummies.iloc[0:300000,0:642]
X_test = df_dummies.iloc[300000:500000,0:642]

In [30]: selector = SelectKBest(f_classif, k=10)
selected_features = selector.fit_transform(X_train, y_train)

scores = selector.scores_

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\feature_selection\_univariate_selection.py:114: UserWarning: Features (36, 38, 350, 381, 410, 528, 544, 552) are constant.
warnings.warn("Features %s are constant." % constant_features_idx,
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\feature_selection\_univariate_selection.py:116: RuntimeWarning: invalid value encountered in true_divide
f = mnb / msw

In [8]: plt.plot(scores)
plt.show()
```

```
In [9]: f_score_indexes = (-selector.scores_).argsort()[1:50]
f_score_indexes

Out[9]: array([631, 636, 629, 627, 639, 625, 624, 641, 618, 619, 11, 12, 637,
        640, 635, 5, 43, 622, 623, 21, 275, 24, 6, 8, 64, 1,
        186, 292, 67, 164, 3, 40, 28, 2, 19, 20, 635, 296, 68,
        260, 355, 42, 368, 188, 237, 13, 30, 170, 208, 171], dtype=int64)

In [90]: scores[464]

Out[90]: 0.005542254653787618

In [13]: df_dummies.head()

Out[13]:
```

	cont0	cont1	cont2	cont3	cont4	cont5	cont6	cont7	cont8	cont9	...	cat16_C	cat16_D	cat17_A	...
0	0.629858	0.855349	0.759439	0.795549	0.881917	0.621672	0.592184	0.791921	0.815254	0.965006	...	0	1	0	0
1	0.370727	0.328929	0.386385	0.541366	0.388982	0.357778	0.600044	0.408701	0.399353	0.927406	...	0	0	1	0
2	0.502272	0.322749	0.343255	0.616352	0.793687	0.552877	0.352113	0.388835	0.412303	0.292696	...	0	1	0	0
3	0.934242	0.707663	0.831147	0.807807	0.800032	0.619147	0.221789	0.897617	0.633669	0.760318	...	0	1	0	0
4	0.254427	0.274514	0.338818	0.277308	0.610578	0.128291	0.578764	0.279167	0.351103	0.357084	...	0	0	0	0

5 rows × 642 columns

```
In [31]: cols = df_dummies.columns

for j in range(642):
    if scores[j] < 3:
        df_dummies = df_dummies.drop(cols[j], axis=1)

In [32]: df_dummies.shape

Out[32]: (500000, 502)

In [33]: X_train = df_dummies.iloc[0:300000,0:502]
X_test = df_dummies.iloc[300000:500000,0:502]

In [37]: X_train.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300000 entries, 0 to 299999
Data columns (total 22 columns):
#   Column      Non-Null Count  Dtype
---  --
0   cat0_A      300000 non-null  uint8
1   cat0_B      300000 non-null  uint8
2   cat10_BU    300000 non-null  uint8
3   cat10_BW    300000 non-null  uint8
4   cat10_CA    300000 non-null  uint8
5   cat10_DG    300000 non-null  uint8
6   cat10_EJ    300000 non-null  uint8
7   cat10_JM    300000 non-null  uint8
8   cat10_KB    300000 non-null  uint8
9   cat10_KM    300000 non-null  uint8
10  cat11_A     300000 non-null  uint8
11  cat11_B     300000 non-null  uint8
12  cat14_A     300000 non-null  uint8
13  cat14_B     300000 non-null  uint8
14  cat15_B     300000 non-null  uint8
15  cat15_D     300000 non-null  uint8
16  cat16_B     300000 non-null  uint8
17  cat16_D     300000 non-null  uint8
18  cat17_D     300000 non-null  uint8
19  cat18_B     300000 non-null  uint8
20  cat18_C     300000 non-null  uint8
21  cat18_D     300000 non-null  uint8
dtypes: uint8(22)
memory usage: 6.3 MB

In [34]: X_train_df, X_test_df, y_train_df, y_test_df = train_test_split(X_train, y_train, test_size=0.4)

In [35]: X_train_df.shape

Out[35]: (180000, 502)
```

## Model Selection

```
In [98]: MLA = [
    ensemble.AdaBoostClassifier(),
    ensemble.BaggingClassifier(),
    ensemble.ExtraTreesClassifier(),
    ensemble.GradientBoostingClassifier(),
    ensemble.RandomForestClassifier(),

    linear_model.LogisticRegressionCV(),
    linear_model.PassiveAggressiveClassifier(),
    linear_model.RidgeClassifierCV(),
    linear_model.SGDClassifier(),
    linear_model.Perceptron(),

    naive_bayes.BernoulliNB(),
    naive_bayes.GaussianNB(),

    neighbors.KNeighborsClassifier(),

    svm.SVC(probability=True),
    svm.NuSVC(probability=True),
    svm.LinearSVC(),

    tree.DecisionTreeClassifier(),
    tree.ExtraTreeClassifier(),

    discriminant_analysis.LinearDiscriminantAnalysis(),
    discriminant_analysis.QuadraticDiscriminantAnalysis(),

    #XGBClassifier()
]

In [48]: model = Adamodel.fit(X_train,y_train)

In [56]: formod = forest.fit(X_train, y_train)

In [23]: logmod = log.fit(X_train, y_train)

In [39]: nbmod = nb.fit(X_train, y_train)
```

## Model Testing

```
In [99]: def MLA_test(X_initial, y_initial, f, t, verbose=0):
    X = X_initial[fit]
    y = y_initial[fit]

    #split dataset in cross-validation with this splitter class: http://scikit-learn.org/stable/module
s/generated/sklearn.model_selection.ShuffleSplit.html#sklearn.model_selection.ShuffleSplit
    #note: this is a non-bias random sample, then t=3 standard deviations (std) from the mean, should
    cv_split = model_selection.ShuffleSplit(n_splits = 10, test_size = .3, train_size = .6, random_stat
e = 0 ) # run model 10x with 60/30 split intentionally leaving out 10%

    #create table to compare MLA metrics
    MLA_columns = ['MLA Name', 'MLA Parameters', 'MLA Train Accuracy Mean', 'MLA Test Accuracy Mean', 'M
LA Test Accuracy %STD', 'MLA Time']
    MLA_compare = pd.DataFrame(columns = MLA_columns)

    #index through MLA and save performance to table
    row_index = 0
    for alg in MLA:
        #set name and parameters
        MLA_name = alg._class_._name_
        MLA_compare.loc[row_index, 'MLA Name'] = MLA_name
        MLA_compare.loc[row_index, 'MLA Parameters'] = str(alg.get_params())

        #score model with cross validation: http://scikit-learn.org/stable/modules/generated/sklearn.mo
del_selection.cross_validate.html#sklearn.model_selection.cross_validate
        cv_results = model_selection.cross_validate(alg, X, y, cv = cv_split, n_jobs=-1, verbose=0, re
turn_train_score=True)

        training_score = cv_results['train_score'].mean()
        test_score = cv_results['test_score'].mean()
        if verbose == 1:
            print('{}\t{}\t{}\t{}\t{}\t{}\t{}\t{}\t{}\t{}\t{}\t{}\t{}\t{}\t{}\t{}\t{}\t{}\t{}\t{}'
                .format(row_index+1, len(MLA)), MLA_name, " - ", training_score, test_score)

        MLA_compare.loc[row_index, 'MLA Time'] = MLA_name
        MLA_compare.loc[row_index, 'MLA Parameters'] = str(alg.get_params())
        MLA_compare.loc[row_index, 'MLA Train Accuracy Mean'] = training_score
        MLA_compare.loc[row_index, 'MLA Test Accuracy Mean'] = test_score
        #if this is a non-bias random sample, then t=3 standard deviations (std) from the mean, should
        statistically capture 99.7% of the subsets
        MLA_compare.loc[row_index, 'MLA Test Accuracy %STD'] = cv_results['test_score'].std()*3 #le
t's know the worst that can happen!

        #save MLA predictions - see section 6 for usage
        #alg.fit(data1[data1_x_bin], df[target])
        #MLA_predict[MLA_name] = alg.predict(df[data1_x_bin])

        row_index+=1

    #print and sort table: https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.sort
_values.html
    MLA_compare.sort_values(by = ['MLA Test Accuracy Mean'], ascending = False, inplace = True)

    return MLA_compare

In [101]: MLA_compare = MLA_test(X_train, y_train, 0, 10000, verbose=1)
MLA_compare
```

	MLA Name	MLA Parameters	MLA Train Accuracy Mean	MLA Test Accuracy Mean	MLA Test Accuracy %STD	MLA Time
4	RandomForestClassifier	{'bootstrap': True, 'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'entropy', 'max_depth': None, 'max_features': 'sqrt', 'max_leaf_nodes': None, 'min_samples_split': 2, 'min_samples_weighted': 1, 'n_estimators': 100, 'oob_score': False, 'random_state': None, 'verbose': 0, 'warm_start': False}	0.99995	0.805933	0.0190252	4.24012
3	GradientBoostingClassifier	{'ccp_alpha': 0.0, 'criterion': 'entropy', 'learning_rate': 0.1, 'max_depth': 3, 'max_features': 'sqrt', 'min_samples_split': 2, 'min_samples_weighted': 1, 'n_estimators': 100, 'oob_score': False, 'random_state': None, 'verbose': 0, 'warm_start': False}	0.813167	0.805633	0.0217368	6.58936
2	ExtraTreesClassifier	{'bootstrap': False, 'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'entropy', 'max_depth': None, 'max_features': 'sqrt', 'max_leaf_nodes': None, 'min_samples_split': 2, 'min_samples_weighted': 1, 'n_estimators': 100, 'oob_score': False, 'random_state': None, 'verbose': 0, 'warm_start': False}	1	0.8038	0.0152525	2.36166
20	XGBClassifier	{'objective': 'binary:logistic', 'base_score': 0.5, 'booster': 'gbtree', 'callbacks': None, 'collect_metrics': True, 'disable_categorical_update': False, 'early_stopping_rounds': None, 'eval_metric': 'logloss', 'feature_select': 'best', 'gamma': 0, 'grow_policy': 'depthwise', 'importance_score': 'weight', 'interaction_constraints': None, 'keep_training_time': False, 'learning_rate': 0.1, 'max_bin': 255, 'max_cat_threshold': 31, 'max_cat_to_onehot': False, 'max_delta_step': 0.05, 'max_depth': 6, 'max_leaf_nodes': None, 'min_child_weight': 1, 'monotone_constraints': None, 'missing': None, 'multi_logit': False, 'n_estimators': 100, 'n_jobs': -1, 'num_parallel_tree': 1, 'one_hot': False, 'open_model': False, 'random_state': None, 'reg_alpha': 0, 'reg_lambda': 0, 'scale_pos_weight': 1, 'seed': 0, 'silent': False, 'sketch_reduce': False, 'subsample': 1, 'tree_method': 'exact', 'verbose': 0, 'watchdog': True}	0.98305	0.798767	0.0102961	4.89539
13	SVC	{'C': 1.0, 'break_ties': False, 'cache_size': 128, 'class_weight': None, 'coef0': 0, 'decision_function_shape': 'raw', 'gamma': 0.5, 'kernel': 'rbf', 'libsvm_verbose': -1, 'max_iter': 1000, 'nu': 0.5, 'probability': True, 'random_state': None, 'shrinking': False, 'tol': 0.001, 'verbose': 0}	0.802833	0.7969	0.0207415	22.9352
0	AdaBoostClassifier	{'algorithm': 'SAMME', 'base_estimator': None, 'n_estimators': 50, 'random_state': None, 'verbose': 0}	0.805467	0.7969	0.0166436	1.87583
5	LogisticRegressionCV	{'Cs': 1.0, 'class_weight': None, 'cv': 'cross_val_score', 'n_jobs': 1, 'penalty': 'l2', 'random_state': None, 'solver': 'lbfgs', 'tol': 0.0001, 'verbose': 0}	0.796117	0.7946	0.0187266	2.0944
15	LinearSVC	{'C': 1.0, 'class_weight': None, 'dual': True, 'fit_intercept': True, 'max_iter': 1000, 'multi_class': 'ovr', 'nu': 0.5, 'penalty': 'l2', 'random_state': None, 'solver': 'libsvm', 'tol': 0.0001, 'verbose': 0}	0.79745	0.794	0.0160873	0.457899
16	LinearDiscriminantAnalysis	{'n_components': None, 'prior': None, 'shrinkage': 'auto', 'solver': 'eigen', 'tol': 0.0001, 'verbose': 0}	0.7956	0.792633	0.0187481	0.0859158
7	RidgeClassifierCV	{'alphas': array([0.1, 1., 10., 100., 1000., 10000., 100000., 1000000., 10000000., 100000000., 1000000000., 10000000000., 100000000000., 1000000000000., 10000000000000., 100000000000000., 1000000000000000., 10000000000000000., 100000000000000000., 1000000000000000000.]), 'cv': 'cross_val_score', 'fit_intercept': True, 'max_iter': 1000, 'multi_class': 'ovr', 'pre_dispatch': 'all', 'random_state': None, 'scoring': 'accuracy', 'solver': 'cholesky', 'tol': 0.0001, 'verbose': 0}	0.793383	0.790867	0.0193142	0.042682
14	NuSVC	{'break_ties': False, 'cache_size': 128, 'class_weight': None, 'coef0': 0, 'decision_function_shape': 'raw', 'gamma': 0.5, 'kernel': 'rbf', 'libsvm_verbose': -1, 'max_iter': 1000, 'nu': 0.5, 'probability': True, 'random_state': None, 'shrinking': False, 'tol': 0.0001, 'verbose': 0}	0.792283	0.787667	0.0209189	25.6482
8	SGDClassifier	{'alpha': 0.0001, 'average': False, 'class_weight': None, 'fit_intercept': True, 'loss': 'hinge', 'max_iter': 1000, 'multi_class': 'ovr', 'n_jobs': 1, 'penalty': 'l2', 'random_state': None, 'shuffle': True, 'tol': 0.0001, 'verbose': 0}	0.7904	0.7869	0.0181467	0.115682
1	BaggingClassifier	{'base_estimator': None, 'bootstrap': True, 'n_estimators': 100, 'random_state': None, 'verbose': 0}	0.982917	0.7838	0.0185213	1.67855
12	KNeighborsClassifier	{'algorithm': 'auto', 'leaf_size': 30, 'metric': 'minkowski', 'metric_p': 2, 'n_neighbors': 5, 'p': 2, 'radius': 1, 'weights': 'distance', 'x_squared_dist': False}	0.838967	0.782533	0.019653	0.281363
10	BernoulliNB	{'alpha': 1.0, 'binarize': 0.0, 'class_prior': None, 'fit_intercept': True, 'force_memory': False, 'irrelevant_features': None, 'laplace': 1.0, 'multinomial': False, 'priors': None, 'smooth': True, 'verbose': 0}	0.73085	0.7278	0.018145	0.020067
9	Perceptron	{'alpha': 0.0001, 'class_weight': None, 'early_stopping': False, 'eta0': 0.01, 'fit_intercept': True, 'max_iter': 1000, 'multi_class': 'ovr', 'penalty': 'l2', 'random_state': None, 'shuffle': True, 'tol': 0.0001, 'verbose': 0}	0.72455	0.727133	0.221373	0.0345262
16	DecisionTreeClassifier	{'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'entropy', 'max_depth': None, 'max_features': 'sqrt', 'max_leaf_nodes': None, 'min_samples_split': 2, 'min_samples_weighted': 1, 'n_estimators': 100, 'oob_score': False, 'random_state': None, 'verbose': 0, 'warm_start': False}	1	0.722767	0.0159502	0.310905
17	ExtraTreeClassifier	{'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'entropy', 'max_depth': None, 'max_features': 'sqrt', 'max_leaf_nodes': None, 'min_samples_split': 2, 'min_samples_weighted': 1, 'n_estimators': 100, 'oob_score': False, 'random_state': None, 'verbose': 0, 'warm_start': False}	1	0.7187	0.0176094	0.0375465
6	PassiveAggressiveClassifier	{'C': 1.0, 'average': False, 'class_weight': None, 'decision_function_shape': 'raw', 'fit_intercept': True, 'gamma': 0.5, 'kernel': 'rbf', 'libsvm_verbose': -1, 'max_iter': 1000, 'nu': 0.5, 'probability': True, 'random_state': None, 'shrinking': False, 'tol': 0.0001, 'verbose': 0}	0.702833	0.7012	0.266824	0.0426862
11	GaussianNB	{'priors': None, 'var_smoothing': 1e-09}	0.6931	0.692733	0.0290923	0.0212302
19	QuadraticDiscriminantAnalysis	{'priors': None, 'reg_param': 0.0, 'store_cova': False, 'tol': 0.0001, 'verbose': 0}	0.665067	0.664867	0.19689	0.0280566

## TensorFlow Model

```
In [36]: df_dummies.shape

Out[36]: (500000, 502)

In [37]: y_train.shape

Out[37]: (300000,)

In [38]: f=0
t = 180000
X = X_train_df[fit]
y = y_train_df[fit]

In [39]: X_train_df.shape

Out[39]: (180000, 502)

In [77]: model = keras.Sequential([
    keras.layers.Dense(60, activation='relu', input_shape= (502,)),
    keras.layers.Dense(40, activation='relu'),
   
```